Homework #3 Programming Assignments
Due: 11:59 pm, October 21, 2018

General instructions

Starting point   Your repository will have now a directory "P3/". Please do not change the name of this repository or the names of any files we have added to it. Please perform a `git pull` to retrieve these files. You will find within it: `boosting.py`, `boosting.sh`, `boosting_check.py`, `boosting_test.py`, `classifier.py`, `data_loader.py`, `decision_stump.py`, `decision_tree.py`, `decision_tree.sh`, `decision_tree_check.py`, `decision_tree_test.py`, `pegasos.py`, `pegasos.sh`, `requirements.txt`

Download the data   Please download **mnist_subset.json** from Piazza/Resource/Homework. Please **DO NOT** push it into your repository when you submit your results. Otherwise, you will get 20% deduction of your score of this assignment.

High-level descriptions

Dataset   We will use **mnist_subset** (images of handwritten digits from 0 to 9). This is the same subset of the full MNIST that we used for Homework 1 and Homework 2. As before, the dataset is stored in a JSON-formated file **mnist_subset.json**. You can access its training, validation, and test splits using the keys 'train', 'valid', and 'test', respectively. For example, suppose we load **mnist_subset.json** to the variable x. Then, $x['train']$ refers to the training set of **mnist_subset**. This set is a list with two elements: $x['train'][0]$ containing the features of size N (samples) $\times D$ (dimension of features), and $x['train'][1]$ containing the corresponding labels of size N.

Tasks

1. You will be asked to implement the linear support vector machine (SVM) for binary classification (Sect. 1). Specifically, you will

    - finish implementing the following three functions—`objective_function`, `pegasos_train`, and `pegasos_test`—in `pegasos.py`. Refer to `pegasos.py` and Sect. 1 for more information.
    - run the script `pegasos.sh` after you finish your implementation. This will output `pegasos.json`.
    - add, commit, and push (1) the `pegasos.py`, and (2) the `pegasos.json` file that you have created.

2. You will be asked to implement the boosting algorithms with decision stump (Sect. 2). Specifically, you will

    - finish implementing the following classes — `boosting`, `decision_stump`, `AdaBoost` and `LogitBoost`. Refer to `boosting.py`, `decision_stump.py` and Sect. 2 for more information.
    - test and debug your code using the provided `boosting_check.py`.
    - run the script `boosting.sh` (which executes `boosting_test.py`). This will output `boosting.json`.
    - add, commit, and push (1) the `boosting.py`, (2) the `decision_stump.py`, and (3) the `boosting.json`.

3. You will be asked to implement the decision tree classifier (Sect. 3). Specifically, you will

- finish implementing the following classes — `DecisionTree`, `TreeNode`. Refer to `decision_tree.py` and Sect. 3 for more information.

- test and debug your code using the provided `decision_tree_check.py`.

- run the script `decision_tree.sh` (which executes `decision_tree_test.py`). This will output `decision_tree.json`.

- add, commit, and push (1) the `decision_tree.py`, (2) the `decision_tree.json`

As in the previous homework, you are not responsible for loading/pre-processing data; we have done that for you (e.g., we have pre-processed the data so it has two classes: digits 0–4 and digits 5–9.). For specific instructions, please refer to text in Sect. 1 and 2 and the instructions in `pegasos.py` and `boosting.py`.

Cautions Please do not import packages that are not listed in the provided code. Follow the instructions in each section strictly to code up your solutions. **Do not change the output format**. **Do not modify the code unless we instruct you to do so**. A homework solution that does not match the provided setup, such as format, name, initializations, etc., **will not** be graded. It is your responsibility to **make sure that your code runs with the provided commands and scripts on the VM**. Finally, make sure that you **git add, commit, and push all the required files**, including your code and generated output files.

Problem 1 Pegasos: a stochastic gradient based solver for linear SVM

In this question, you will build a linear SVM (i.e. without kernel) classifier using the Pegasos algorithm [1]. Given a training set $\{(\mathbf{x}_n \in \mathbb{R}^D, y_n \in \{1, -1\})\}_{n=1}^N$, the primal formulation of linear SVM can be written as L2-regularized hinge loss as shown in the class:

$$\min_{\mathbf{w}} \frac{\lambda}{2} ||\mathbf{w}||_2^2 + \frac{1}{N} \sum_n \max\{0, 1 - y_n \mathbf{w}^T \mathbf{x}_n\}. \tag{1}$$

Note that here we include the bias term b into parameter $\mathbf{w}$ by appending $\mathbf{x}$ with 1, which is slightly different from what was discussed in the class.

Instead of turning it into dual formulation, we are going to solve the primal formulation directly with a gradient-base algorithm. Recall for (batch) gradient descent, at each iteration of parameter update, we compute the gradients for all data points and take the average (or sum). When the training set is large (i.e., N is a large number), this is often too computationally expensive. Stochastic gradient descent with mini-batch alleviates this issue by computing the gradient on a *subset* of the data at each iteration.

Pegasos, a representative solver of eq. (1), is exactly applying stochastic gradient descent with mini-batch to this problem, with an extra step of projection described below (this is also called projected stochastic gradient descent). The pseudocode of Pegasos is given in Algorithm 1. At the t-th iteration of parameter update, we first take a mini-batch of data A_t of size K [step (3)], and define $A_t^+ \subset A_t$ to be the subset of samples for which $\mathbf{w}_t$ suffers a non-zero loss [step (4)]. Next we set the learning rate $\eta_t = 1/(\lambda t)$ [step (5)]. We then perform a **two-step parameter update** as follows. We first compute $(1 - \eta_t \lambda)\mathbf{w}_t$ and for all samples $(\mathbf{x}, y) \in A_t^+$ we add the vector $\frac{y \eta_t}{K}\mathbf{x}$ to $(1 - \eta_t \lambda)\mathbf{w}_t$. We denote the resulting vector by $\mathbf{w}_{t+\frac{1}{2}}$ [step (6)]. This step is in fact exactly stochastic gradient descent: $\mathbf{w}_{t+\frac{1}{2}} = \mathbf{w}_t - \eta_t \nabla_t$ where

$$\nabla_t = \lambda \mathbf{w}_t - \frac{1}{K} \sum_{(\mathbf{x}, y) \in A_t^+} y\mathbf{x}$$

is the (sub)gradient of the objective function on the mini-batch A_t at $\mathbf{w}_t$. Last, we set $\mathbf{w}_{t+1}$ to be the projection of $\mathbf{w}_{t+\frac{1}{2}}$ onto the set

$$B = \{\mathbf{w} : ||\mathbf{w}||_2 \leq 1/\sqrt{\lambda}\}$$

to control its L2 norm. This can be obtained simply by scaling $\mathbf{w}_{t+1}$ by $\min\left\{1, \frac{1/\sqrt{\lambda}}{||\mathbf{w}_{t+\frac{1}{2}}||}\right\}$ [step (e)].

For more details of Pegasos algorithm you may refer to the original paper [1].

Now you are going to implement Pegasos and train a binary linear SVM classifier on the dataset **mnist_subset.json**. You are going to implement three functions—`objective_function`, `pegasos_train`, and `pegasos_test`—in a script named `pegasos.py`. You will find detailed programming instructions in the script.

1.1 Finish the implementation of the function `objective_function`, which corresponds to the objective function in the primal formulation of SVM.

1.2 Finish the implementation of the function `pegasos_train`, which corresponds to Algorithm 1.

Algorithm 1 The Pegasos algorithm

Require: a training set $S = \{(\mathbf{x}_n \in \mathbb{R}^D, y_n \in \{1, -1\})\}_{n=1}^N$, the total number of iterations T, the batch size K, and the regularization parameter λ.

Output: the last weight $\mathbf{w}_{T+1}$.

1: **Initialization:** $\mathbf{w}_1 = \mathbf{0}$
2: **for** $t = 1, 2, \cdots, T$ **do**
3:     Randomly choose a subset of data $A_t \in S$ (with replacement) such that $|A_t| = K$
4:     Set $A_t^+ = \{(\mathbf{x}, y) \in A_t : y\mathbf{w}_t^T\mathbf{x}_t < 1\}$
5:     Set $\eta_t = \dfrac{1}{\lambda t}$
6:     Set $\mathbf{w}_{t+\frac{1}{2}} = (1 - \eta_t\lambda)\mathbf{w}_t + \dfrac{\eta_t}{K}\sum_{(\mathbf{x},y) \in A_t^+} y\mathbf{x}$
7:     Set $\mathbf{w}_{t+1} = \min\left\{1, \dfrac{1/\sqrt{\lambda}}{||\mathbf{w}_{t+\frac{1}{2}}||}\right\}\mathbf{w}_{t+\frac{1}{2}}$
8: **end for**

1.3 After you train your model, run your classifier on the test set and report the accuracy, which is defined as:

$$\frac{\text{\# of correctly classified test samples}}{\text{\# of test samples}}$$

Finish the implementation of the function `pegasos_test`.

1.4 After you complete above steps, run `pegasos.sh`, which will run the Pegasos algorithm for 500 iterations with 6 settings (mini-batch size $K = 100$ with different $\lambda \in \{0.01, 0.1, 1\}$ and $\lambda = 0.1$ with different $K \in \{1, 10, 1000\}$), and output a `pegasos.json` that records the test accuracy and the value of objective function at each iteration during the training process.

What to do and submit: run script `pegasos.sh`. It will take `mnist_subset.json` as input and generate `pegasos.json`. Add, commit, and push both `pegasos.py` and `pegasos.json` before the due date.

1.5 Based on `pegasos.json`, you are encouraged to make plots of the objective function value versus the number of iterations (i.e., a convergence curve) in different settings of λ and K, but you are not required to submit these plots.

References

[1] Shai Shalev-Shwartz, Yoram Singer, Nathan Srebro, Andrew Cotter *Mathematical Programming, 2011, Volume 127, Number 1, Page 3*. Pegasos: primal estimated sub-gradient solver for SVM.

Problem 2 Boosting

In the lectures, we discussed boosting algorithms, which construct a strong (binary) classifier based on iteratively adding one weak (binary) classifier into it. A weak classifier is learned to (approximately) maximize the weighted training accuracy at the corresponding iteration, and the weight of each training example is updated after every iteration.

In this question, you will implement decision stumps as weak classifiers and AdaBoost. For simplicity, instead of implementing an actual base algorithm that learns decision stumps, we fix a finite set of decision stumps $\mathcal{H}$ in advance and pick the one with smallest weighted error at each round of boosting.

2.1 General boosting framework A boosting algorithm sequentially learns β_t and $h_t \in \mathcal{H}$ for $t = 0, \ldots, T - 1$ and finally constructs a strong classifier as $H(\mathbf{x}) = \text{sign}\left[\sum_{t=0}^{T-1} \beta_t h_t(\mathbf{x})\right]$.

Please read the class "Boosting" defined in `boosting.py`, and then complete the class by implementing the "TODO" part(s) as indicated in `boosting.py`.

2.2 Decision stump A decision stump $h \in \mathcal{H}$ is a classifier characterized by a triplet $(s \in \{+1, -1\}, b \in \mathbb{R}, d \in \{0, 1, \cdots, D - 1\})$ such that

$$h_{(s,b,d)}(\mathbf{x}) = \begin{cases} s, & \text{if } x_d > b, \\ -s, & \text{otherwise.} \end{cases} \tag{2}$$

That is, each decision stump function only looks at a single dimension x_d of the input vector $\mathbf{x}$, and check whether x_d is larger than b. Then s decides which label to predict if $x_d > b$.

Please first read `classifier.py` and `decision_stump.py`, and then complete the class "DecisionStump" by implementing the "TODO" part(s) as indicated in `decision_stump.py`.

2.3 AdaBoost AdaBoost is a powerful and popular boosting method, summarized in Algorithm 2.

Algorithm 2 The AdaBoost algorithm

Require: $\mathcal{H}$: A set of classifiers, where $h \in \mathcal{H}$ takes a D-dimensional vector as input and outputs either $+1$ or -1, a training set $\{(\mathbf{x}_n \in \mathbb{R}^D, y_n \in \{+1, -1\})\}_{n=0}^{N-1}$, the total number of iterations T.

Ensure: Learn $H(\mathbf{x}) = \text{sign}\left[\sum_{t=0}^{T-1} \beta_t h_t(\mathbf{x})\right]$, where $h_t \in \mathcal{H}$ and $\beta_t \in \mathbb{R}$.

1: **Initialization** $D_0(n) = \frac{1}{N}, \forall n \in \{0, 1, \cdots, N - 1\}$.
2: **for** $t = 0, 1, \cdots, T - 1$ **do**
3:     find $h_t = \text{argmin}_{h \in \mathcal{H}} \sum_n D_t(n) \mathbb{I}[y_n \neq h(\mathbf{x}_n)]$
4:     compute $\epsilon_t = \sum_n D_t(n) \mathbb{I}[y_n \neq h_t(\mathbf{x}_n)]$
5:     compute $\beta_t = \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t}$
6:     compute $D_{t+1}(n) = \begin{cases} D_t(n) \exp(-\beta_t), & \text{if } y_n = h_t(\mathbf{x}_n) \\ D_t(n) \exp(\beta_t), & \text{otherwise} \end{cases}, \forall n \in \{0, 1, \cdots, N - 1\}$
7:     normalize $D_{t+1}(n) \leftarrow \frac{D_{t+1}(n)}{\sum_{n'} D_{t+1}(n')}, \forall n \in \{0, 1, \cdots, N - 1\}$
8: **end for**

Please read the class "AdaBoost" defined in `boosting.py`, and then complete the class by implementing the "TODO" part(s) as indicated in `boosting.py`.

2.4 Final submission Please run `boosting.sh`, which will generate `boosting.json`. Add, commit, and push `boosting.py`, `decision_stump.py` and `boosting.json` before the due date.

Problem 3 Decision Tree

In this question, you will implement a simple decision tree algorithm. For simplicity, only discrete features are considered here.

3.1 Conditional Entropy Recall that the quality of a split can be measured by the conditional entropy. Please read `decision_tree.py` and complete the function "conditional_entropy" (where we use base 2 for logarithm).

3.2 Tree construction The building of a decision tree involves growing and pruning. For simplicity, in this programming set you are only asked to grow a tree.

As discussed a tree can be constructed by recursively splitting the nodes if needed, as summarized in Algorithm 3 (whose interface is slightly different from the pseudocode discussed in the class). Please read `decision_tree.py` and complete the function "split".

Algorithm 3 The recursive procedure of decision tree algorithm

---

1: **function** TREENODE.SPLIT(self)
2:     find the best attribute to split the node
3:     split the node into child_nodes
4:     **for** child_node in child_nodes **do**
5:         **if** child_node.splittable **then**   ▷ See Slide 27 of Lec 7 for the three "non-splittable" cases
6:             child_node.split()
7:         **end if**
8:     **end for**
9: **end function**

3.3 Final submission Please run `decision_tree.sh`, which will generate `decision_tree.json`. Add, commit, and push `decision_tree.py`, `decision_tree.json` before the due date.